

# Data Cleaning

Andrew Jaffe

January 6, 2016

# Data

- ▶ We will be using multiple data sets in this lecture:
  - ▶ Salary, Monument, Circulator, and Restaurant from OpenBaltimore: <https://data.baltimorecity.gov/browse?limitTo=datasets>
  - ▶ Gap Minder - very interesting way of viewing longitudinal data
    - ▶ Data is here - <http://www.gapminder.org/data/>
  - ▶ [http://spreadsheets.google.com/pub?key=rMsQHawT0bBb6\\_U2ESjKXYw&output=xls](http://spreadsheets.google.com/pub?key=rMsQHawT0bBb6_U2ESjKXYw&output=xls)

# Data Cleaning

In general, data cleaning is a process of investigating your data for inaccuracies, or recoding it in a way that makes it more manageable.

**MOST IMPORTANT RULE - LOOK AT YOUR DATA!**

Again - `table`, `summarize`, `is.na`, `any`, `all` are useful.

# Data Cleaning

```
table(c(0, 1, 2, 3, NA, 3, 3, 2,2, 3),  
      useNA="ifany")
```

```
##  
##      0      1      2      3 <NA>  
##      1      1      3      4      1
```

```
table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),  
      useNA="always")
```

```
##  
##      0      1      2      3 <NA>  
##      1      1      4      4      0
```

```
tab <- table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),  
             c(0, 1, 2, 3, 2, 3, 3, 4, 4, 3),  
             useNA="always")  
margin.table(tab, 2)
```

```
prop.table(tab)
```

```
##
##           0    1    2    3    4 <NA>
##  0      0.1 0.0 0.0 0.0 0.0  0.0
##  1      0.0 0.1 0.0 0.0 0.0  0.0
##  2      0.0 0.0 0.2 0.0 0.2  0.0
##  3      0.0 0.0 0.0 0.4 0.0  0.0
## <NA> 0.0 0.0 0.0 0.0 0.0  0.0
```

```
prop.table(tab,1)
```

```
##
##           0    1    2    3    4 <NA>
##  0      1.0 0.0 0.0 0.0 0.0  0.0
##  1      0.0 1.0 0.0 0.0 0.0  0.0
##  2      0.0 0.0 0.5 0.0 0.5  0.0
##  3      0.0 0.0 0.0 1.0 0.0  0.0
## <NA>
```

# Download Salary FY2014 Data

[https://data.baltimorecity.gov/City-Government/  
Baltimore-City-Employee-Salaries-FY2014/2j28-xzd7](https://data.baltimorecity.gov/City-Government/Baltimore-City-Employee-Salaries-FY2014/2j28-xzd7)

Download as a CSV and then read it into R as the variable Sal

```
Sal = read.csv("http://www.aejaffe.com/winterR_2016/data/Ba  
as.is = TRUE)
```

# Data Cleaning

- ▶ `any()` - checks if there are any TRUES
- ▶ `all()` - checks if ALL are true

```
Sal[1:4,]
```

##	Name	JobTitle	Agency
## 1	Aaron,Keontae E	AIDE BLUE CHIP	W02
## 2	Aaron,Patricia G	Facilities/Office Services II	A03
## 3	Aaron,Petra L	ASSISTANT STATE'S ATTORNEY	A29
## 4	Abaineh,Yohannes T	EPIDEMIOLOGIST	A65

  

##	Agency	HireDate	AnnualSalary	Gross
## 1	Youth Summer	06/10/2013	\$11310.00	\$87
## 2	OED-Employment Dev	10/24/1979	\$53428.00	\$5286
## 3	States Attorneys Office	09/25/2006	\$68300.00	\$6743
## 4	HLTH-Health Department	07/23/2009	\$62000.00	\$5865

```
any(is.na(Sal$Name))
```

## Example of Cleaning:

For example, let's say gender was coded as Male, M, m, Female, F, f. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

In R, you can simply do something like:

```
data$gender[data$gender %in%  
  c("Male", "M", "m")] <- "Male"
```



Sometimes though, it's not so simple. That's where functions that find patterns come in very useful.

```
table(gender)
```

```
## gender
##      F FeMAle FEMALE      Fm      M      Ma      mAle      Male
##      75      82      74      89      89      79      87      89
##      Man   Woman
##      73      80
```

# Find/Replace and Regular Expressions

- ▶ R can do much more than find exact matches for a whole string
- ▶ Like Perl and other languages, it can use regular expressions.
- ▶ What are regular expressions?
- ▶ Ways to search for specific strings
- ▶ Can be very complicated or simple
- ▶ Highly Useful

## 'Find' functions

`grep`: `grep`, `grep1`, `regexpr` and `gregexpr` search for matches to argument `pattern` within each element of a character vector: they differ in the format of and amount of detail in the results.

`grep(pattern, x, fixed=FALSE)`, where:

- ▶ `pattern` = character string containing a regular expression to be matched in the given character vector.
- ▶ `x` = a character vector where matches are sought, or an object which can be coerced by `as.character` to a character vector.
- ▶ If `fixed=TRUE`, it will do exact matching for the phrase anywhere in the vector (regular find)

## 'Find' functions

```
grep("Rawlings",Sal$Name)
```

```
## [1] 13832 13833 13834 13835
```

These are the indices/elements where the pattern match occurs

grep() returns something similar to which() on a logical statement

## 'Find' functions

```
grep("Rawlings",Sal$Name)
```

```
## [1] 13832 13833 13834 13835
```

```
grep("Rawlings",Sal$Name,value=TRUE)
```

```
## [1] "Rawlings,Kellye A"           "Rawlings,MarqWell D"  
## [3] "Rawlings,Paula M"           "Rawlings-Blake,Stephanie C"
```

```
Sal[grep("Rawlings",Sal$Name),]
```

##	Name	JobTitle	Agency	HireDate	AnnualSalary	GrossPay
## 13832	Rawlings,Kellye A	EMERGENCY DISPATCHER	M-B Info Technology	01/06/2003	\$47980.00	\$6840.00
## 13833	Rawlings,MarqWell D	AIDE BLUE CHIP				
## 13834	Rawlings,Paula M	COMMUNITY AIDE				
## 13835	Rawlings-Blake,Stephanie C	MAYOR				

## grep() Options

```
head(grep("Tajhgh",Sal$Name, value=TRUE))
```

```
## [1] "Reynold,Tajhgh J"
```

```
grep("Jaffe",Sal$Name)
```

```
## [1] 8603
```

```
length(grep("Jaffe",Sal$Name))
```

```
## [1] 1
```

# A bit on Regular Expressions

- ▶ `http:`  
`//www.regular-expressions.info/reference.html`
- ▶ They can use to match a large number of strings in one statement
- ▶ `.` matches any single character
- ▶ `*` means repeat as many (even if 0) more times the last character
- ▶ `?` makes the last thing optional

# Using Regular Expressions

- ▶ Look for any name that starts with:
  - ▶ Payne at the beginning,
  - ▶ Leonard and then an S
  - ▶ Spence then a capital C

```
grep("Payne.*", x=Sal$Name, value=TRUE)
```

```
## [1] "Payne El,Jackie"      "Payne Johnson,Nickole A"
## [3] "Payne,Chanel"        "Payne,Connie T"
## [5] "Payne,Denise I"      "Payne,Dominic R"
## [7] "Payne,James R"       "Payne,Jasman T"
## [9] "Payne,Joey D"        "Payne,Jordan A"
## [11] "Payne,Karen V"       "Payne,Karen V"
## [13] "Payne,Leonard S"    "Payne,Mary A"
## [15] "Payne,Micah W"       "Payne,Michael C"
## [17] "Payne,Michael N"     "Payne,Morag"
## [19] "Payne,Nora M"        "Payne,Shelley F"
```



```
grep("Leonard.?S", x=Sal$Name, value=TRUE)
```

```
## [1] "Payne,Leonard S"      "Szumlanski,Leonard S"
```

```
grep("Spence.*C.*", x=Sal$Name, value=TRUE)
```

```
## [1] "Greene,Spencer C"      "Spencer,Charles A"    "Spencer
```

```
## [4] "Spencer,Clarence W"    "Spencer,Michael C"
```

# Replace

Let's say we wanted to sort the data set by Annual Salary:

```
class(Sal$AnnualSalary)
```

```
## [1] "character"
```

```
sort(c("1", "2", "10")) # not sort correctly (order simply)
```

```
## [1] "1" "10" "2"
```

```
order(c("1", "2", "10"))
```

```
## [1] 1 3 2
```

# Replace

So we must change the annual pay into a numeric:

```
head(as.numeric(Sal$AnnualSalary), 4)
```

```
## Warning in head(as.numeric(Sal$AnnualSalary), 4): NAs in  
## coercion
```

```
## [1] NA NA NA NA
```

R didn't like the \$ so it thought turned them all to NA.

sub() and gsub() can do the replacing part.

## Replacing and subbing

Now we can replace the \$ with nothing (used fixed=TRUE because \$ means something in regular expressions):

```
Sal$AnnualSalary <- as.numeric(gsub(pattern="$", replacement=" ",
                                     Sal$AnnualSalary, fixed=TRUE))
Sal <- Sal[order(Sal$AnnualSalary, decreasing=TRUE), ] # use order
Sal[1:5, c("Name", "AnnualSalary", "JobTitle")]
```

##		Name	AnnualSalary	JobTitle
## 1222	Bernstein, Gregg L	238772	STATE'S ATTORNEY	
## 3175	Charles, Ronnie E	200000	EXECUTIVE LEVEL 1	
## 985	Batts, Anthony W	193800	EXECUTIVE LEVEL 1	
## 1343	Black, Harry E	190000	EXECUTIVE LEVEL 1	
## 16352	Swift, Michael	187200	CONTRACT SERV SPEC	

# Useful String Functions

## Useful String functions

- ▶ `toupper()`, `tolower()` - uppercase or lowercase your data:
- ▶ `str_trim()` (in the `stringr` package) - will trim whitespace
- ▶ `nchar` - get the number of characters in a string
- ▶ `substr(x, start, stop)` - substrings from position start to position stop
- ▶ `strsplit(x, split)` - splits strings up - returns list!
- ▶ `paste()` - paste strings together - look at ?paste

## Paste

Paste can be very useful for joining vectors together:

```
paste("Visit", 1:5, sep="_")
```

```
## [1] "Visit_1" "Visit_2" "Visit_3" "Visit_4" "Visit_5"
```

```
paste("Visit", 1:5, sep="_", collapse=" ")
```

```
## [1] "Visit_1 Visit_2 Visit_3 Visit_4 Visit_5"
```

```
paste("To", "is going be the ", "we go to the store!", sep=
```

```
## [1] "Today is going be the day we go to the store!"
```

*# and paste0 can be even simpler see ?paste0*

```
paste0("Visit",1:5)
```

```
## [1] "Visit1" "Visit2" "Visit3" "Visit4" "Visit5"
```

```
paste(1:5, letters[1:5], sep="_")
```

```
## [1] "1_a" "2_b" "3_c" "4_d" "5_e"
```

```
paste(6:10, 11:15, 2000:2005, sep="/")
```

```
## [1] "6/11/2000" "7/12/2001" "8/13/2002" "9/14/2003"  
## [6] "6/11/2005"
```

```
paste(paste("x", 1:5, sep=""), collapse="+")
```

```
## [1] "x1+x2+x3+x4+x5"
```

# Strsplit

```
x <- c("I really", "like writing", "R code")  
y <- strsplit(x, split=" ")  
y[[2]]
```

```
## [1] "like"      "writing"
```

```
sapply(y, "[", 1) # on the fly
```

```
## [1] "I"        "like" "R"
```

```
sapply(y, "[", 2) # on the fly
```

```
## [1] "really" "writing" "code"
```



# Data Merging/Append

- ▶ Merging - joining data sets together - usually on key variables, usually “id”
- ▶ `merge()` is the most common way to do this with data sets
- ▶ `rbind/cbind` - row/column bind, respectively
  - ▶ `rbind` is the equivalent of “appending” in Stata or “setting” in SAS
  - ▶ `cbind` allows you to add columns in addition to the previous ways
- ▶ `reshape2` package also has a lot of information about different ways to reshape data (wide to long, etc) - but has a different (and sometimes more intuitive syntax)
- ▶ `t()` is a function that will transpose the data

# Merging

```
base <- data.frame(id=1:10, Age= seq(55,60, length=10))  
base[1:2,]
```

```
##   id      Age  
## 1  1 55.00000  
## 2  2 55.55556
```

```
visits <- data.frame(id=rep(1:8, 3), visit= rep(1:3, 8),  
                     Outcome= seq(10,50, length=24))  
visits[1:2,]
```

```
##   id visit Outcome  
## 1  1     1 10.00000  
## 2  2     2 11.73913
```

```
merged.data <- merge(base, visits, by="id")
merged.data[1:5,]
```

```
##      id      Age visit Outcome
## 1  1 55.00000      1 10.00000
## 2  1 55.00000      3 23.91304
## 3  1 55.00000      2 37.82609
## 4  2 55.55556      2 11.73913
## 5  2 55.55556      1 25.65217
```

```
dim(merged.data)
```

```
## [1] 24  4
```

```
all.data <- merge(base, visits, by="id", all=TRUE)
tail(all.data)
```

##		id	Age	visit	Outcome
##	21	7	58.33333	2	48.26087
##	22	8	58.88889	2	22.17391
##	23	8	58.88889	1	36.08696
##	24	8	58.88889	3	50.00000
##	25	9	59.44444	NA	NA
##	26	10	60.00000	NA	NA

```
dim(all.data)
```

```
## [1] 26 4
```

## Aside: Dates

You can convert date-like strings in the Date class

(<http://www.statmethods.net/input/dates.html> for more info)

```
circ = read.csv("http://www.aejaffe.com/winterR_2016/data/CircData.csv",
               as.is=TRUE)
head(sort(circ$date))
```

```
## [1] "01/01/2011" "01/01/2012" "01/01/2013" "01/02/2011"
## [6] "01/02/2013"
```

```
circ$date <- as.Date(circ$date, "%m/%d/%Y") # creating a date
head(circ$date)
```

```
## [1] "2010-01-11" "2010-01-12" "2010-01-13" "2010-01-14"
## [6] "2010-01-16"
```

```
head(sort(circ$date))
```

# Data Reshaping

Disclaimer: the `reshape` command in R is not remarkably intuitive.

- ▶ Wide - multiple measurements are variables / columns so that the data gets wider with more measurements
- ▶ Long - multiple measurements are rows so data gets longer with more measurements
- ▶ One example would be many ids with multiple visits

## Example of Long/Wide

```
head(wide)
```

```
##   id visit1 visit2 visit3  
## 1  1   Good   Good   Bad
```

```
head(long)
```

```
##   id visit Outcome  
## 1  1     1    Good  
## 2  1     2    Good  
## 3  1     3    Bad
```

# Data Reshaping

- ▶ Good resource:

<http://www.ats.ucla.edu/stat/r/faq/reshape.htm>

```
head(Indometh) # this is long
```

```
##      Subject time conc
## 1          1 0.25 1.50
## 2          1 0.50 0.94
## 3          1 0.75 0.78
## 4          1 1.00 0.48
## 5          1 1.25 0.37
## 6          1 2.00 0.19
```



# Data Reshaping

```
wide <- reshape(Indometh, v.names = "conc", idvar = "Subject",  
                 timevar = "time", direction = "wide")  
head(wide)
```

##	Subject	conc.0.25	conc.0.5	conc.0.75	conc.1	conc.1.25
## 1	1	1.50	0.94	0.78	0.48	0.37
## 12	2	2.03	1.63	0.71	0.70	0.64
## 23	3	2.72	1.49	1.16	0.80	0.80
## 34	4	1.85	1.39	1.02	0.89	0.55
## 45	5	2.05	1.04	0.81	0.39	0.30
## 56	6	2.31	1.44	1.03	0.84	0.64

  

##	conc.4	conc.5	conc.6	conc.8
## 1	0.11	0.08	0.07	0.05
## 12	0.20	0.25	0.12	0.08
## 23	0.12	0.11	0.08	0.08
## 34	0.11	0.10	0.07	0.07
## 45	0.11	0.08	0.10	0.06
## 56	0.17	0.13	0.10	0.09

# Data Reshaping

```
dim(Indometh)
```

```
## [1] 66  3
```

```
wide
```

```
##      Subject conc.0.25 conc.0.5 conc.0.75 conc.1 conc.1.25
```

```
## 1          1        1.50        0.94        0.78        0.48        0.37
```

```
## 12         2        2.03        1.63        0.71        0.70        0.64
```

```
## 23         3        2.72        1.49        1.16        0.80        0.80
```

```
## 34         4        1.85        1.39        1.02        0.89        0.59
```

```
## 45         5        2.05        1.04        0.81        0.39        0.30
```

```
## 56         6        2.31        1.44        1.03        0.84        0.64
```

```
##      conc.4 conc.5 conc.6 conc.8
```

```
## 1      0.11  0.08  0.07  0.05
```

```
## 12     0.20  0.25  0.12  0.08
```

```
## 23     0.12  0.11  0.08  0.08
```

```
## 34     0.11  0.10  0.07  0.07
```

# Data Reshaping

- ▶ If you've reshaped a data set - to get it back, just reshape it again

```
reshape(wide, direction = "long")[1:10,]
```

##	Subject	time	conc
## 1.0.25	1	0.25	1.50
## 2.0.25	2	0.25	2.03
## 3.0.25	3	0.25	2.72
## 4.0.25	4	0.25	1.85
## 5.0.25	5	0.25	2.05
## 6.0.25	6	0.25	2.31
## 1.0.5	1	0.50	0.94
## 2.0.5	2	0.50	1.63
## 3.0.5	3	0.50	1.49
## 4.0.5	4	0.50	1.39

Note the row name change

## Data Reshaping - A Better Example

```
## Loading required package: rJava
## Loading required package: xlsxjars
```

```
destfile = tempfile(fileext = ".xlsx")
download.file("http://www.aejaffe.com/winterR_2016/data/inc
TB <- read.xlsx(file = destfile,
                sheetName = "Data")
head(TB, 1)
```

```
## TB.incidence..all.forms..per.100.000.population.per.ye
## 1 Afghanistan
## X1992 X1993 X1994 X1995 X1996 X1997 X1998 X1999 X2000
## 1 168 168 168 168 168 168 168 168 168
## X2004 X2005 X2006 X2007 NA.
## 1 168 168 168 168 NA
```

```
TB$NA. <- NULL
head(TB, 1)
```

# Data Reshaping - A Better Example

```
colnames(TB) <- c("Country", paste("Year",  
                                     1990:2007, sep="."))  
head(TB,1)
```

```
##      Country Year.1990 Year.1991 Year.1992 Year.1993 Year.1994  
## 1 Afghanistan      168      168      168      168      168  
##      Year.1996 Year.1997 Year.1998 Year.1999 Year.2000 Year.2001  
## 1      168      168      168      168      168      168  
##      Year.2003 Year.2004 Year.2005 Year.2006 Year.2007 Year.2008  
## 1      168      168      168      168      168      168
```

## Data Reshaping - More is better!

```
TB.long <- reshape(TB, idvar="Country",  
                    v.names="Cases", times=1990:2007,  
                    direction="long", timevar="Year",  
                    varying = paste("Year", 1990:2007, sep=""),  
                    sep="")  
head(TB.long, 4)
```

```
##              Country Year Cases  
## Afghanistan.1990    Afghanistan 1990    168  
## Albania.1990         Albania 1990     25  
## Algeria.1990         Algeria 1990     38  
## American Samoa.1990 American Samoa 1990    21
```

```
rownames(TB.long) <- NULL  
head(TB.long, 4)
```

```
##              Country Year Cases  
## 1    Afghanistan 1990    168  
## 2    Albania 1990     25
```